RAG Technical Notes

Retrieval-Augmented Generation improves LLM reliability.

1. Chunking: split long documents into small chunks.
2. Embedding: convert chunks into vectors.
3. Retrieval: find the most relevant chunks.
4. Generation: answer based on retrieved context.

Medical RAG helps diagnosis and clinical decision support.